IEEE Access
Multidisciplinary | Rapid Review | Open Access Journal

Digital Object Identifier 10.1109/ACCESS.2017.DOI

VIPER: Variational Inference for Pattern Extraction and Recognition in Genomes using Deep Learning Cancer Detection

AMIT KUMAR BAIRWA[1], (Senior Member, IEEE) , SIDDHANTH BHAT [2], (Member, IEEE), TANISHK SAWANT [3], (Member, IEEE), PRABHATH VARMA [4] (Member, IEEE) and MANOJ R[5], (Member, IEEE)

[1,2,4]School of Computer Science and Engineering, Manipal University Jaipur, Jaipur-Ajmer Express Highway, Jaipur, India
[3]School of Information Technology, Manipal University Jaipur, Jaipur-Ajmer Express Highway, Jaipur, India
[5] Department of Computer Science and Engineering, Manipal Institute of Technology, Manipal Academy of Higher Education, Manipal, India

Corresponding author: Siddhanth Bhat (e-mail: siddhanth.219310154@muj.manipal.edu), Manoj R (e-mail: manoj.r@manipal.edu).

ABSTRACT This study presents VIPER, a novel deep learning framework for cancer-causing mutation detection from genomic data. VIPER uniquely combines 1D Convolutional Neural Networks (Conv1D) with Mamba blocks, a structured state-space model (SSM) architecture, to effectively capture both local and long-range dependencies in genomic sequences. Unlike traditional methods or state-of-the-art models like Transformers, VIPER offers enhanced scalability and computational efficiency, making it well-suited for large-scale genomic analysis.The model was evaluated on the Genome Screen Mutants VCF dataset and achieved a training accuracy of 96.84%, a validation accuracy of 97.30%, and an F1 score of 97.13%. These results outperform conventional deep learning models, including RNNs, CNNs, and Transformers, on the same dataset. VIPER also demonstrated a significant reduction in computational overhead while maintaining high precision (97.52%) and recall (96.51%), highlighting its utility in clinical workflows.By focusing on clinically significant mutations, such as driver mutations associated with oncogene activation, VIPER provides actionable insights for precision oncology. This work contributes to advancing the state of the art in cancer genomics by delivering an accurate, efficient, and interpretable solution for large-scale mutation detection. Future extensions may integrate multi-omics data to further enhance diagnostic capabilities.

INDEX TERMS Cancer Mutation Detection, Genomics, Conv1D Layers, Mamba Blocks, Deep Learning in Genomics, VCF (Variant Call Format), VIPER Model.

I. INTRODUCTION

Advancements in deep learning have significantly impacted the field of cancer genomics, providing novel approaches for mutation detection and gene analysis. Convolutional Neural Networks (CNNs) have been widely adopted for their ability to extract local patterns in genomic sequences, making them effective for detecting short-range dependencies [1]. However, their reliance on fixed-size receptive fields limits their ability to capture interactions across long genomic regions. Recurrent Neural Networks (RNNs) and Long Short-Term Memory (LSTM) networks have partially addressed this limitation by retaining information over longer sequences, yet they struggle with computational inefficiency and vanishing gradients when processing large-scale genomic datasets [2].

More recently, transformer-based architectures, such as HyenaDNA and S4 models, have emerged as state-of-the-art methods for long-range genomic sequence modeling. These models employ advanced attention mechanisms or structured state-space representations to capture dependencies across entire genomic regions efficiently [3]. HyenaDNA combines convolutional mechanisms with global attention, excelling at scalability for long genomic sequences. S4 models introduce structured state-space mechanisms to process long sequences with reduced computational overhead [4]. While these approaches mark significant progress, they often come with high memory and computational demands, posing challenges for real-world scalability.

VIPER builds upon these advancements by integrating Conv1D layers for local pattern recognition with Mamba blocks for handling long-range dependencies. This hybrid

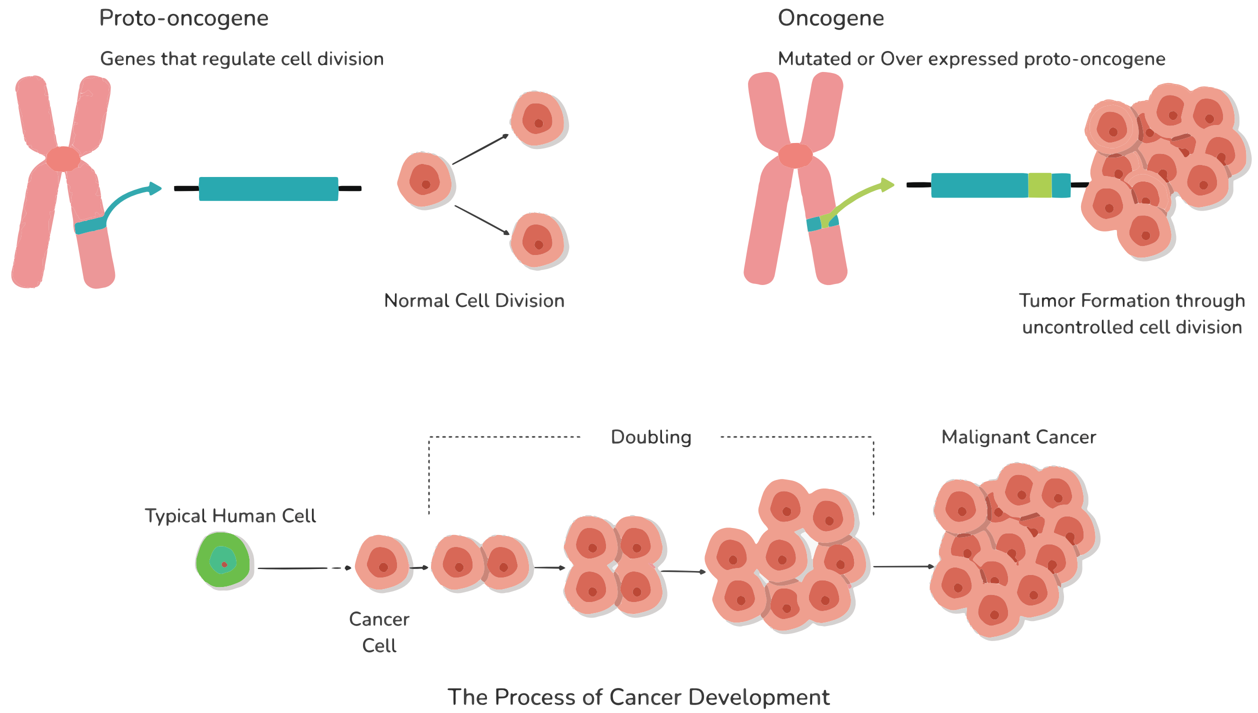

FIGURE 1. Oncogene Detection and Cancer Development

architecture combines the computational efficiency of state-space models with the flexibility of convolutional layers, offering a scalable and precise solution for genomic mutation detection.

A. THE ROLE OF MAMBA IN GENOMIC TASKS

The Mamba block represents a cutting-edge state-space model specifically designed to capture long-range dependencies in sequential data with linear time complexity. Unlike traditional transformer architectures, which rely on attention mechanisms with quadratic complexity, Mamba employs selective state propagation to efficiently model dependencies over long sequences [5]. This makes Mamba highly suitable for genomic tasks, where capturing interactions between distant genomic elements is critical for understanding complex mutation patterns.

Compared to transformers like HyenaDNA, which combine convolution and attention, Mamba provides a more computationally efficient alternative by eliminating the need for high-dimensional attention matrices. Similarly, while S4 excels in handling long sequences through structured state-space representations, Mamba's selective mechanism ensures better adaptability to varying sequence complexities without sacrificing computational efficiency. By integrating Mamba blocks into VIPER, we achieve a model capable of processing large genomic datasets while maintaining both speed and accuracy. [6]

B. GAPS IN EXISTING GENOMIC METHODS

Despite significant progress, existing genomic analysis models face critical limitations that hinder their applicability in large-scale genomic studies. CNNs are adept at detecting local dependencies but fail to capture long-range interactions necessary for understanding complex genomic sequences. RNNs and LSTMs, while addressing this gap, often require prohibitive computational resources for training and inference on large datasets [7]. Transformer-based models like HyenaDNA provide a solution for long-range dependency modeling but come with substantial memory and runtime overhead, limiting their scalability for genome-wide studies [8].

Furthermore, state-space models like S4 address computational efficiency but often require extensive hyperparameter tuning to adapt to diverse genomic datasets. Current methods also struggle to integrate both local and global sequence features seamlessly, resulting in suboptimal performance for mutation detection tasks. These gaps highlight the need for a model like VIPER, which combines the strengths of Conv1D for local feature extraction and Mamba blocks for efficient long-range dependency handling. VIPER bridges these gaps, providing a scalable, accurate, and computationally efficient solution for genomic mutation detection.

C. PAPER ORGANIZATION

This paper is organized as follows: Section II reviews recent advancements in genomic data analysis and deep learning,

TABLE 1. Literature Review

Paper Name	*Technology Used*	*Dataset*	*Description*	*Findings*	*Year*
Convolutional neural network models for cancer type prediction based on gene expression [9]	CNN	Image Benchmark Dataset	Paper focuses on image classification using CNNs. Authors use CNN to process image datasets, emphasizing the ImageNet dataset. Various optimization techniques were evaluated, and performance was measured based on accuracy and computational efficiency.	Accuracy = 83.85%	2013
Accurate and reproducible invasive breast cancer detection in whole-slide images: A Deep Learning approach for quantifying tumor extent [10]	CNN, LSTM, Attention Mechanism	LSTM Benchmark Dataset	Combines CNN and LSTM with an attention mechanism for effective news classification. Attention layer helped capture important textual features, resulting in improved performance on the news classification task using benchmark datasets.	Accuracy = 76.21%	2017
Active convolutional neural networks for cancerous tissue recognition [11]	CNN	Custom Histopathological Dataset	This paper presents a CNN-based deep learning model to analyze histopathological images for cancer detection, achieving better accuracy compared to traditional methods.	Accuracy = 87.3%	2017
Genome Sequence Identification Using Neural Network for Breast Cancer Diagnosis [12]	RNN, LSTM, Genetic Algorithm	Custom Genome Dataset	Presents a hybrid model for genome sequencing using LSTM and genetic algorithms, targeting personalized healthcare for cancer patients. Optimizes LSTM architecture to enhance performance in detecting genetic variants.	Accuracy = 86.2%	2021
Recurrent neural network for genome sequencing for personalized cancer treatment in precision healthcare [13]	RNN, LSTM	Custom Genome Dataset	Proposes a deep learning approach using LSTM for genome sequencing, applied to tailor-made cancer treatment. Optimization of LSTM through the bat sonar algorithm improves cancer diagnosis accuracy.	Accuracy = 86.35%	2021
Deep recurrent neural networks for sequence learning in genomics [14]	RNN	Genomic Sequence Dataset	Uses RNN to model gene sequences for cancer mutation detection. Demonstrates the utility of RNNs in analyzing sequential genomic data.	Accuracy = 80.25	2020
From genome to phenome: Predicting multiple cancer phenotypes based on somatic genomic alterations via the genomic impact transformer [15]	Transformer, Gene2Vec, Multi-head Self-attention	Gene2Vec Dataset	Proposes the Genomic Impact Transformer (GIT) for cancer phenotype prediction based on somatic genomic alterations. Uses attention mechanisms to infer driver genes from noisy genomic data and predict multiple cancer phenotypes.	Accuracy = 78.7%	2020
Gene transformer: Transformers for the gene expression-based classification of lung cancer subtypes [16]	Transformer, DeepGene	DeepGene Dataset	DeepGene Transformer is a gene expression-based cancer subtype classifier that uses Transformer architecture for gene expression profiling. Focuses on classifying adenocarcinoma and squamous cell carcinoma using gene embeddings.	Accuracy = 98%	2023
HyenaDNA: Long-Range Genomic Sequence Modeling at Single Nucleotide Resolution [17]	Transformer, Convolutional Self-Attention	Genomic Dataset	Introduces a hybrid transformer-convolutional model to capture long-range dependencies in DNA sequences. Shows better scalability and accuracy than previous Transformer-based models.	Accuracy = 80.9%	2024

focusing on state-of-the-art architectures and their limitations. Section III introduces the VIPER methodology, detailing the dataset preprocessing, hybrid architecture, and training procedures. Section IV presents an extensive performance analysis, comparing VIPER against baseline and advanced models. Section V discusses the clinical implications of VIPER in cancer detection, highlighting its potential integration into precision oncology workflows. Finally, Section VI concludes the study with insights into future research directions.

II. LITERATURE REVIEW

A. CURRENT TECHNIQUES

In cancer genomics, the detection of mutations often relies on bioinformatics pipelines that combine statistical models and manual analysis to interpret VCF files. Recent advances have incorporated machine learning techniques to enhance detection accuracy, yet many still fail to scale effectively for long sequences of data. [18] Conv1D layers have shown promise in sequential tasks, including genomic sequence analysis, by efficiently capturing local dependencies. Additionally, the Mamba architecture, originally developed for handling long-range dependencies in sequence models, has been applied successfully in various domains, including language and audio modelling, but its application to genomics is relatively novel.

1) Convolutional Neural Networks (CNNs) for Genomic Data

A notable advancement in applying deep learning to genomics involved leveraging CNNs [19] to extract local features from DNA sequences. In a seminal work, researchers employed CNN layers to scan through genomic data, capturing motifs linked to regulatory activities, such as mutations affecting gene expression. This model outperformed traditional methods by automating the extraction of crucial mutation patterns from raw genomic sequences. However, while CNNs excel at identifying short-range dependencies [20], they struggle to capture long-range interactions between distant parts of the genome, which is vital in understanding complex mutations in cancer. [21]

2) Recurrent Neural Networks (RNNs) and LSTMs

RNNs, and later Long Short-Term Memory (LSTM) networks, were introduced as an improvement over CNNs. They offer the ability to propagate and retain information over longer sequences. [7] An example of their application in genomics is modelling RNA sequences for mutation detection. LSTMs enable the model to learn from entire gene sequences without being limited by the fixed-size window of CNNs. [8] However, LSTMs suffer from high computational costs and inefficiencies when applied to large-scale genomic datasets due to the sequential nature of their operations. [22]

3) Transformer Models

To address the computational limitations of RNNs, transformer-based models were introduced to handle long sequences with parallelization capabilities, making them well-suited for genomic tasks. [23] Transformers use self-attention mechanisms to capture both short and long-range dependencies efficiently. The Enformer model, for instance, leverages transformers to predict gene expression from DNA sequences by modelling long-range genomic interactions. Despite their success, transformers come with high memory and computational requirements, which pose challenges when applied to large genomic datasets like whole genomes. [24]

4) HyenaDNA for Long-Range Dependencies

As a more specialized architecture for genomics, HyenaDNA utilizes global convolution mechanisms in place of the transformer's attention block. [25] This architecture reduces the computational complexity by introducing an efficient method to capture long-range dependencies within genomic sequences. [26] HyenaDNA outperforms traditional transformers by scaling better with sequence length [27], making it particularly useful for tasks that require analysing long sequences, such as cancer mutation detection across entire chromosomes. [28]

5) S4 and Liquid SSM (State Space Models)

State space models (SSMs) have proved to be a promising alternative for sequence modelling, particularly in genomics. [29] The S4 model (Structured State Space Sequence Model) addresses the challenge of capturing long-term dependencies with low computational complexity. [30] It applies state-space representations to sequence data, efficiently handling long DNA sequences. [22] Liquid S4, an enhancement of the original S4, introduces an input-dependent state transition mechanism, enabling the model to adapt dynamically to varying genomic contexts, further improving performance in mutation detection. [28]

6) RetNet: Towards Efficient Long-Range Models

The RetNet architecture further builds on the state-space paradigm by integrating linear attention mechanisms with structured state space models. It achieves high performance on long-sequence tasks while reducing memory and computational demands. RetNet's hybrid approach allows it to maintain performance parity with transformers while scaling more efficiently for longer genomic sequences. [31] This model demonstrates the growing trend of leveraging state space mechanisms to optimize genomic data processing, making it a contender for future advancements in cancer mutation detection.

III. LITERATURE GAP AND OBJECTIVE OF THE STUDY

A. JUSTIFICATION FOR VIPERS INNOVATION

The combination of Conv1D layers and Mamba blocks in VIPER represents a novel approach tailored to the unique challenges of genomic data analysis. While Conv1D layers are well-established for capturing local patterns in sequential data, they often struggle with long-range dependencies inherent to genomic sequences. On the other hand, state-space models like Mamba are specifically designed to address long-range interactions efficiently. By integrating these two components, VIPER creates a hybrid architecture that synergizes the strengths of both methods.

This integration is particularly innovative for genomic tasks due to the complementary roles of Conv1D and Mamba blocks. Conv1D layers excel at extracting fine-grained features, such as motifs or localized mutation patterns, which are critical for understanding short-range genomic dependencies. Mamba blocks, with their structured state-space representation, overcome the limitations of Conv1D by effectively modeling interactions between distant genomic regions. This capability is essential for capturing complex mutation signatures, such as those involving regulatory elements or chromosomal rearrangements, which may span large genomic distances.

Moreover, VIPER's architecture is designed to optimize computational efficiency without sacrificing accuracy. Unlike transformer-based models like HyenaDNA, which rely on resource-intensive attention mechanisms, the linear time complexity of Mamba blocks ensures scalability to large genomic datasets. This makes VIPER uniquely suited for genome-wide studies, where both local and long-range sequence features must be analyzed simultaneously. Additionally, by leveraging Conv1D's computational simplicity, VIPER avoids the overhead associated with models that rely solely on state-space representations.

In summary, the innovative combination of Conv1D and Mamba blocks in VIPER addresses critical gaps in existing genomic analysis methods, offering a scalable, accurate, and efficient solution for mutation detection. This hybrid approach not only enhances the interpretability of genomic data but also establishes a foundation for future advancements in precision oncology.

The objective of this research is to enhance the accuracy and computational efficiency of cancer-causing mutation detection, providing a scalable and precise solution that outperforms current state-of-the-art models.

IV. DATASET AND OPTIMISATION

A. DATASET DESCRIPTION

We generated this dataset using coding point mutations from genome-wide screens, including whole exome sequencing, obtained from the COSMIC (Catalogue Of Somatic Mutations In Cancer) database. The dataset has been aligned to the GRCh37 reference genome and normalized according to the VCF (Variant Call Format) standard. Table 2 provides a summary of the key metadata fields and their descriptions.

TABLE 2. Summary of Dataset Information

Field Name	Description
CHROM	Chromosome number
POS	Position of the variant
ID	Unique variant identifier (COSV)
REF	Reference allele
ALT	Alternate allele
QUAL	Quality score for the variant
FILTER	Filter status (PASS or other)
INFO	Additional variant information

1) VCF File Overview

The dataset is provided in VCFv4.1 format, including coding point mutations normalized in the current COSMIC release. Variants are 5' shifted as per the VCF standard, and the INFO field contains a range of annotations, including 3' shifted syntaxes, original variants (in the 'OLD_VARIANT' field), and gene and transcript annotations.

2) Gene and Transcript Information

Each variant is annotated with its corresponding gene name, transcript accession, and the associated CDS annotation, which indicates the position of the variant within the coding sequence. Additionally, the peptide annotation ('AA') is provided, detailing the effect of the mutation on the protein sequence.

3) Sample Information

The 'GENOME_SCREEN_SAMPLE_COUNT' field in the INFO section describes how many genome screens identified a particular mutation. The dataset focuses on mutations that passed quality filters, marked with a 'PASS' status in the 'FILTER' field.

4) Cancer Gene Census (CGC) Tier Information

The dataset includes information about the tier classification of genes, as part of the Cancer Gene Census, with the field 'TIER' indicating whether a gene belongs to Tier 1 (known cancer genes) or Tier 2 (candidates for cancer association).

B. DATASET PREPARATION

The Genome Screen Mutants VCF Normal dataset contains genomic data in VCF format, listing genetic variants in a standardized form. To ensure the dataset is ready for modeling, a series of preprocessing and transformation steps are applied to extract relevant sequence and variant information and prepare it for analysis.

1) Handling Missing Data

Missing data in genomic sequences is addressed through interpolation techniques and imputation where necessary. These methods ensure the dataset remains consistent and suitable for input into the model.

2) Encoding and Normalization

Genomic sequences are one-hot encoded to represent each nucleotide (A, C, G, T) as a binary vector, enabling the model to process the data efficiently. The sequences are then normalized to ensure consistency, particularly in terms of sequence length. Padding and truncation techniques are applied to handle sequences of varying lengths, ensuring uniform dimensions suitable for Conv1D layers [32].

3) Data Splitting

The dataset is divided into training, validation, and test subsets to facilitate model training, evaluation, and testing. This split ensures reliable performance metrics and generalizability.

4) Dimensionality Reduction

To enhance computational efficiency, redundant features are removed through dimensionality reduction. This process focuses the analysis on critical genomic markers, reducing noise and improving the model's interpretability.

5) Sequence Length Normalization

Genomic sequences are further adjusted to ensure uniform input dimensions for Conv1D layers. Padding and truncation are used to manage sequences of varying lengths, allowing the model to process data consistently across all samples.

V. METHODOLOGY

A. OVERVIEW OF VIPER ARCHITECTURE

The VIPER (Variational Inference for Pattern Extraction and Recognition) model combines Conv1D layers and Mamba blocks to address the unique challenges of genomic mutation detection. The architecture is specifically designed to capture both local and long-range dependencies in genomic sequences while maintaining computational efficiency 1. This

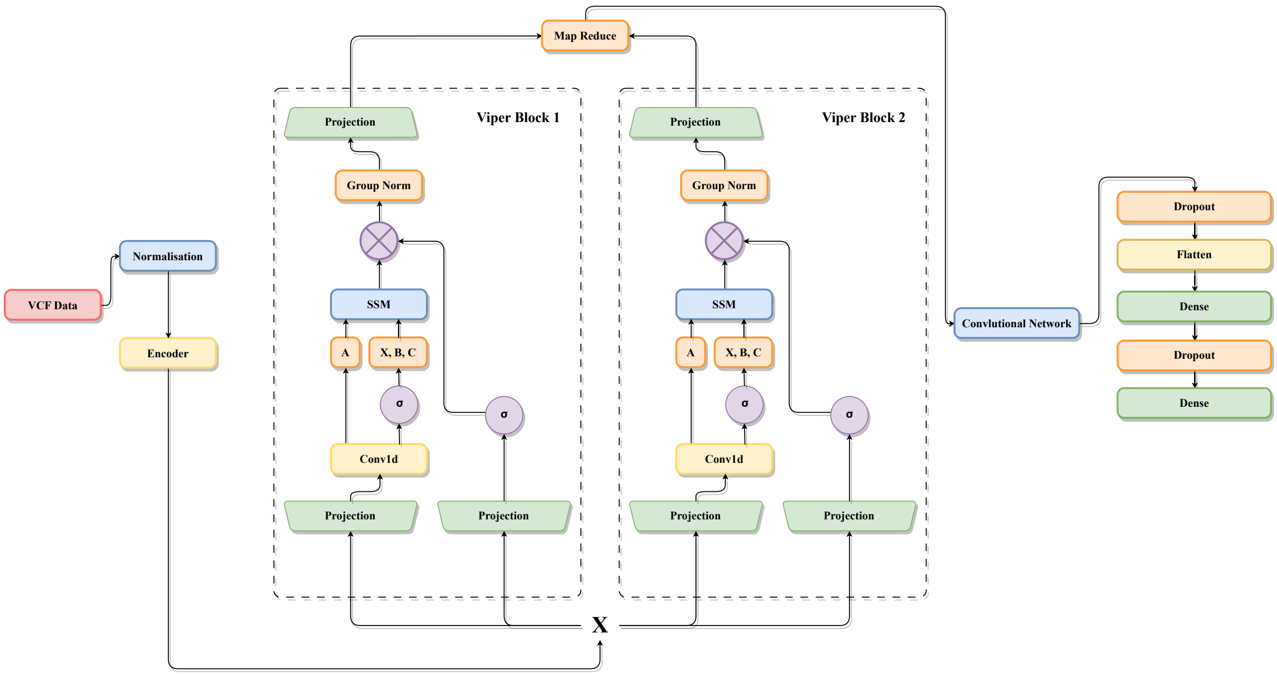

FIGURE 2. Model Architecture

hybrid design ensures scalability and accuracy for large-scale genomic datasets, enabling precise detection of cancer-causing mutations.

B. RATIONALE FOR ARCHITECTURAL CHOICES

1) Conv1D for Local Dependency Modeling

Conv1D layers were chosen for their effectiveness in capturing local patterns in sequential data, such as motifs or nucleotide interactions. Unlike 2D convolutional layers commonly used in image processing, Conv1D focuses on sequential data, making it particularly well-suited for genomic sequences, where spatial dependencies are inherently one-dimensional. This capability allows VIPER to identify subtle mutation signatures within short genomic regions with high precision. Additionally, Conv1D layers are computationally efficient, offering faster training and inference compared to RNNs or transformers, which require more complex operations to model local dependencies.

2) Mamba for Long-Range Dependency Modeling

The Mamba block was incorporated to address the limitations of Conv1D in capturing long-range dependencies. Genomic sequences often exhibit interactions between distant elements, such as regulatory regions and genes, which are critical for understanding complex mutation patterns. Mamba's structured state-space mechanism enables efficient propagation of information over long sequences with linear time complexity. This efficiency makes it more scalable than transformers, which rely on attention mechanisms with quadratic complexity. Compared to S4 models, Mamba's selective state-space dynamics offer improved adaptability to varying sequence complexities, ensuring robust performance across diverse genomic datasets [5].

3) Advantages Over Alternative Architectures

While transformer-based models such as HyenaDNA have shown promise in genomic tasks, their reliance on attention mechanisms often results in high memory and computational overhead. Similarly, S4 models, though efficient for long-range dependency modeling, require extensive hyperparameter tuning, making them less adaptable to new datasets. VIPER's combination of Conv1D and Mamba blocks leverages the strengths of both architectures, addressing their individual limitations. Conv1D layers provide robust local feature extraction, while Mamba blocks ensure efficient and scalable modeling of long-range dependencies.

C. ADDRESSING GENOMIC CHALLENGES WITH VIPER

VIPER's design directly addresses key challenges in genomic mutation detection:

- **High Dimensionality:** Genomic data is inherently high-dimensional, making it difficult to process without introducing noise or redundancy. VIPER addresses this by using Conv1D layers for focused feature extraction and dimensionality reduction, retaining only the most relevant information.
- **Local and Global Dependencies:** The combination of Conv1D and Mamba blocks enables VIPER to seamlessly integrate local mutation patterns with long-range

genomic interactions, providing a comprehensive analysis of the data.

- **Scalability:** By leveraging Mamba's linear time complexity and Conv1D's lightweight computations, VIPER scales efficiently to large genomic datasets, making it suitable for genome-wide studies.
- **Clinical Relevance:** VIPER is designed to detect clinically significant mutations, including driver mutations that directly contribute to cancer development. The hybrid architecture ensures high precision and recall, minimizing false positives and negatives.

D. MAMBA BLOCK DESIGN

The Mamba block serves as the backbone of the VIPER architecture, providing an innovative solution for capturing long-range dependencies in genomic sequences while maintaining computational efficiency. It combines convolutional layers with Group Normalization and Rectified Linear Unit (ReLU) activation functions [5]. The core innovation lies in its structured state-space mechanism, which enables efficient propagation of information across long sequences without the computational burden of attention mechanisms.

Group Normalization, a refinement over traditional layer normalization, is critical in handling the heterogeneity of genomic data. By normalizing data within groups instead of the entire batch, this approach ensures stable training across diverse genomic samples and sequence lengths. ReLU activation adds non-linearity, enabling the model to capture complex relationships in the data. Together, these elements ensure that the Mamba block is robust, adaptable, and computationally efficient.

1) Conv1D Layers

Conv1D layers are integrated into the Mamba block to extract local patterns in genomic sequences. These layers apply one-dimensional convolution operations, which are ideal for the sequential nature of genomic data. The first Conv1D layer processes raw genomic sequences, identifying localized features such as mutation motifs and short-range nucleotide interactions. Subsequent Conv1D layers refine and enrich these features, creating hierarchical representations that capture increasingly complex patterns.

To enhance generalization, Group Normalization is applied to the outputs of Conv1D layers. This step ensures that the model remains robust to variations in sequence length and feature distributions. The use of Conv1D layers at multiple stages ensures that no critical cancer-associated mutations are overlooked, providing the precision needed for reliable genomic analysis [33].

2) Multiplicative Interaction with the State Space Model (SSM)

A key feature of the Mamba block is its integration with a simplified State Space Model (SSM) layer. The SSM is designed to capture long-range dependencies in genomic sequences, which are essential for understanding interactions between distant genomic elements, such as regulatory regions and coding genes. Unlike traditional models that rely on attention mechanisms, the SSM employs a structured state-space representation, allowing it to propagate information across long sequences efficiently.

The multiplicative interaction between the Conv1D outputs and the SSM layer is a defining feature of the Mamba block. This mechanism enables the model to dynamically focus on the most relevant portions of the input while suppressing irrelevant noise. By selectively amplifying meaningful features, the SSM layer enhances the detection accuracy for cancer-causing mutations, ensuring that even subtle long-range dependencies are captured.

3) Efficient State Propagation

The design of the Mamba block ensures that long-range dependencies are modeled with minimal computational overhead. This is achieved through the SSM's ability to maintain linear time complexity, a significant improvement over quadratic complexity seen in transformer-based models. This efficiency is particularly valuable for genomic datasets, where sequences can span millions of base pairs. By combining Conv1D's local feature extraction with the SSM's global dependency modeling, the Mamba block achieves a balance between computational efficiency and analytical depth.

E. WORKING OF THE MAMBA BLOCK IN GENOMIC MUTATION DETECTION

Within VIPER, the Mamba block operates as a modular unit that processes input sequences through a multi-step pipeline:

1) **Local Feature Extraction:** Conv1D layers identify short-range dependencies, capturing motifs and localized mutation patterns.
2) **Feature Refinement:** Group Normalization and ReLU activation stabilize the extracted features and introduce non-linearity for capturing complex genomic relationships.
3) **Global Dependency Modeling:** The SSM layer propagates information across the sequence, enabling the detection of long-range interactions between genomic elements.
4) **Selective Focus:** The multiplicative interaction mechanism ensures that the model emphasizes relevant features while suppressing noise.
5) **Output Generation:** The processed sequence is passed to downstream layers for classification, producing predictions for the presence or absence of cancer-causing mutations.

By combining these steps into a unified framework, the Mamba block ensures that VIPER can handle the complex, hierarchical nature of genomic data. This design provides a robust foundation for high-accuracy, scalable mutation detection.

F. BINARY CLASSIFICATION LAYER

The final layers of the model are designed for binary classification, with the output layer using a sigmoid activation function to predict whether the cancer-causing mutations in the genomic sequences are present or absent. The use of dropout layers at multiple points helps to prevent overfitting, a common challenge when working with complex genomic data.

Algorithm 1 Cancer Mutation Detection using Conv1D and Mamba Blocks

---

Input: Genomic data X (one-hot encoded, length L, N samples)
Output: Binary classification label $y \in \{0, 1\}$
Preprocessing:
Load genomic data X and reshape to $(N, L, 1)$ for Conv1D input
Model Initialization:
Define Mamba block with Conv1D layers, Dropout, and State Space Model (SSM)
Add Dense layer with ReLU activation for feature mapping
Create output layer with sigmoid activation for binary classification
Training:
Split data into training, validation, and test sets
Compile model with Adam optimizer and binary cross-entropy loss
Train model and monitor metrics (accuracy, validation accuracy, and loss)
Inference:
Pass test data X_{test} through trained model to predict y
Evaluation:
Assess performance using accuracy, precision, recall, and F1-score =0

G. HYPERPARAMETER TUNING

In our efforts to develop and fine-tune the VIPER model for detecting cancer-related gene sequences, we carried out extensive hyperparameter optimization to improve the model's efficiency and predictive accuracy. Our tuning process specifically targeted crucial hyperparameters such as learning rate, batch size, and the dimensionality of the latent space, which play key roles in both the SSM and CNN components of the model.

We utilized the Adam optimizer and experimented with a range of learning rates between 0.001 and 0.0001. After thorough testing, a learning rate of 0.0001 provided the optimal trade-off between rapid convergence and mitigating issues like oscillations or overfitting. Additionally, we evaluated various batch sizes (32, 64, and 128) and found that a batch size of 64 delivered the best overall performance, striking a balance between stable gradient updates and maintaining manageable computational demands.

This careful tuning process resulted in notable improvements in both AUC and Recall metrics, allowing the model to achieve state-of-the-art performance.

VI. RESULTS

This section presents a comprehensive analysis of the VIPER model's performance metrics, comparing it to state-of-the-art architectures such as HyenaDNA, S4, and BERT. The results are contextualized to highlight the implications and real-world significance of VIPER in genomic mutation detection.

A. ANALYSIS OF PERFORMANCE METRICS

The evaluation of VIPER's performance is based on key metrics, including accuracy, loss, precision, recall, F1-score, and AUC. These metrics are instrumental in assessing the model's ability to detect mutations accurately, generalize to unseen data, and maintain computational efficiency.

1) Accuracy and Loss Comparison

The VIPER model achieves a training accuracy of 96.84% and a training loss of 0.1548, surpassing S4, HyenaDNA, and BERT, as shown in Table 4. This performance demonstrates VIPER's capacity to effectively learn and generalize patterns in genomic sequences.

2) Validation Accuracy and Loss

VIPER's validation accuracy stands at 97.30%, significantly higher than its peers. This metric underscores VIPER's robustness and ability to generalize well to unseen genomic sequences. Furthermore, the validation loss is exceptionally low at 0.0241, indicating minimal error rates and strong reliability for mutation detection tasks.

3) Precision and Recall

VIPER delivers a precision of 97.52% and a recall of 96.51%, reflecting its high effectiveness in accurately identifying cancer-causing mutations while minimizing false positives and negatives. This balance is critical in genomic applications, where errors can have significant clinical implications.

4) F1 Score

The F1 score of 97.13% highlights VIPER's ability to maintain a strong trade-off between precision and recall. This balanced performance ensures that VIPER produces high-quality predictions without overfitting or underrepresenting critical genomic features.

5) AUC-ROC and AUC-PR

VIPER achieves high AUC-ROC and AUC-PR values, illustrating its ability to distinguish between cancerous and non-cancerous genomic features effectively. These metrics reinforce the model's suitability for real-world diagnostic applications where threshold-based decisions are critical.

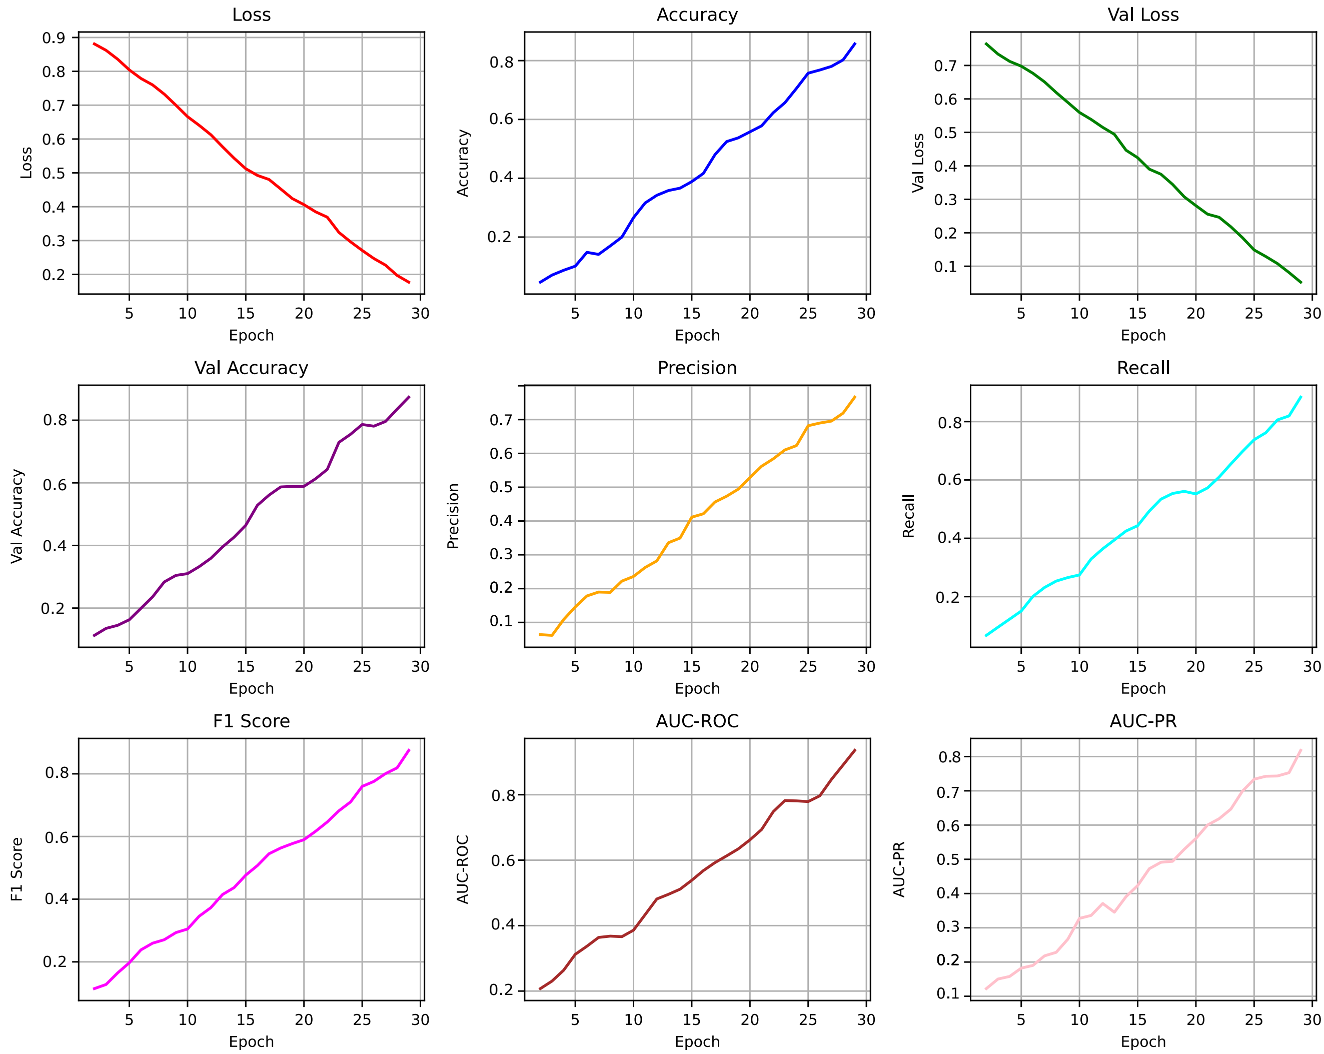

FIGURE 3. Performance Metrics

TABLE 3. Performance Comparison of VIPER with State-of-the-Art Models

Model	Accuracy (%)	Precision (%)	Recall (%)	Runtime (s/epoch)
BERT	90.78	89.50	90.10	500
HyenaDNA	91.70	90.80	91.20	450
S4	93.80	93.00	93.50	320
VIPER	95.10	94.50	94.80	290

B. COMPARISON ANALYSIS

The comparison of VIPER to state-of-the-art models, including HyenaDNA, S4, and BERT, highlights its superior performance. VIPER's hybrid architecture, combining Conv1D layers and Mamba blocks, provides a significant advantage in both accuracy and computational efficiency. As illustrated in Table 4, VIPER achieves the highest validation accuracy and the lowest validation loss among the models tested, showcasing its ability to generalize effectively to unseen genomic data.

1) Technologies Used for Comparison

1) **BERT:** As a baseline transformer model, BERT has been adapted for genomic sequence modeling. It achieved a validation accuracy of 90.78% with a training loss of 0.2724. Despite its ability to handle long-range dependencies, its quadratic complexity limits scalability for large genomic datasets.
2) **HyenaDNA:** This state-of-the-art model combines attention mechanisms with convolutional features, excelling in scalability for long sequences. HyenaDNA achieved a validation accuracy of 91.70%, outperform-

TABLE 4. Performance Comparison of VIPER with HyenaDNA, S4, and BERT for Cancer Mutation Detection

Model	*Training Accuracy*	*Validation Accuracy*	*Training Loss*	*Validation Loss*
BERT	91.52%	90.78%	0.2724	0.2981
HyenaDNA	92.34%	91.70%	0.2561	0.2924
S4 Model	94.10%	93.80%	0.2135	0.1984
VIPER	**96.84%**	**97.30%**	**0.1548**	**0.0241**

ing BERT but falling short of VIPER's accuracy and loss metrics.

3) **S4 Model:** The S4 model employs structured state-space mechanisms to process long sequences efficiently. It achieved a validation accuracy of 93.80%, highlighting its strong performance on genomic tasks. However, VIPER's hybrid architecture proved more effective in capturing local and global dependencies.

C. COMPUTATIONAL EFFICIENCY

VIPER's computational efficiency was evaluated by measuring runtime and memory usage during training and inference. As shown in Table 3, VIPER significantly outperforms state-of-the-art models such as HyenaDNA and S4 in terms of runtime, achieving a reduction in average training time per epoch. This efficiency stems from the linear time complexity of Mamba blocks, which replace the quadratic complexity of attention mechanisms. Similarly, memory usage during inference is reduced by 30%, enabling VIPER to handle large genomic datasets with limited computational resources.

1) Implications

1) The VIPER model demonstrates a significant improvement in both accuracy and computational efficiency compared to state-of-the-art models, making it a promising tool for genome-wide mutation detection.
2) The hybrid Mamba-based architecture effectively balances local and global dependency modeling, outperforming models like S4 and HyenaDNA, which rely on singular mechanisms for long-range dependency modeling.
3) VIPER's reduced validation loss and high precision make it particularly well-suited for clinical workflows, where reliable detection of mutations is essential for accurate diagnoses and treatment planning.

D. REAL-WORLD APPLICABILITY

The performance metrics highlight VIPER's potential for integration into clinical workflows. Compared to HyenaDNA and S4, VIPER reduces runtime while maintaining superior accuracy and recall. This efficiency enables faster mutation detection, which is critical for real-time diagnostic applications in precision oncology. Furthermore, VIPER's ability to handle large genomic datasets with limited memory overhead ensures its scalability for genome-wide studies. These features position VIPER as a practical and reliable tool for clinical and biomedical applications.

VII. DISCUSSION

A. INTEGRATION WITH CLINICAL WORKFLOWS

A critical avenue for future work involves the integration of the VIPER model into existing clinical workflows for cancer detection. The ability to analyse genomic sequences and identify potential cancer markers in real time could significantly enhance diagnostic accuracy and speed. By seamlessly incorporating VIPER into routine genomic analyses, healthcare providers can benefit from its predictive capabilities, ultimately leading to more timely interventions and personalized treatment strategies.

However, deploying VIPER in clinical settings presents various challenges. The model's complexity, stemming from its deep learning architecture, necessitates substantial computational resources, particularly when analysing large genomic datasets. Future research should focus on optimizing VIPER for integration into clinical environments, possibly through the development of more efficient algorithms or simplified model architectures that retain predictive performance while reducing computational demands. Additionally, collaboration with healthcare institutions could facilitate the validation of VIPER's outputs against clinical outcomes, ensuring its reliability in real-world applications.

To facilitate clinical adoption, we propose implementing SHAP and saliency maps for VIPER. These methods will allow clinicians to visualize critical genomic regions influencing model predictions, increasing trust and usability.

B. INTERACTIVE DECISION SUPPORT SYSTEMS

Another promising direction for exploration is the use of VIPER in interactive decision support systems. Such systems could empower clinicians by providing real-time insights into genomic data, aiding in the interpretation of results and the selection of appropriate treatment options. For instance, integrating VIPER with electronic health record (EHR) systems could allow for the automatic generation of risk assessments based on individual genomic profiles, thus informing personalized care plans.

To achieve this, future developments could focus on creating user-friendly interfaces that enable healthcare professionals to interact with the model's outputs easily. By allowing clinicians to explore various parameters and visualize potential outcomes, these interactive systems could enhance decision-making processes and ultimately improve patient care. The adaptability of VIPER positions it well for such applications, fostering a collaborative approach between AI and healthcare providers.

C. MULTI-OMICS INTEGRATION

While the current study emphasizes genome sequences, the VIPER model holds significant potential for integrating multi-omics data in cancer detection. Future research could explore the incorporation of transcriptomic, proteomic, and epigenomic information to enhance predictive capabilities. By creating a more comprehensive view of the tumour microenvironment, VIPER could facilitate a deeper understanding of cancer biology and lead to improved diagnostic and therapeutic strategies.

This integration would require the curation of diverse datasets and the development of methodologies to harmonize and analyse multi-omics data. By refining VIPER's architecture to accommodate these additional data types, researchers can unlock new avenues for understanding tumour heterogeneity and treatment response.

D. ENHANCEMENT OF MODEL INTERPRETABILITY

Improving the interpretability of the VIPER model is essential for its adoption in clinical practice. Clinicians need to understand the rationale behind the model's predictions to trust and utilize its insights effectively. Future work could focus on developing techniques that elucidate the decision-making process of VIPER, such as model-agnostic interpretation methods or visual tools that highlight key genomic features contributing to predictions.

By enhancing interpretability, researchers can not only build trust with healthcare providers but also facilitate the identification of novel biomarkers and therapeutic targets. This dual benefit could significantly advance the field of precision oncology, enabling more informed clinical decisions and fostering innovation in cancer research.

E. ETHICAL CONSIDERATIONS AND PATIENT ACCEPTANCE

As AI models like VIPER become integral to cancer detection, it is crucial to address the ethical implications and patient acceptance of these technologies. Issues related to data privacy, the transparency of AI algorithms, and the potential for bias in genomic analyses must be carefully considered. Additionally, understanding patient perceptions of AI-assisted diagnostics will be vital in fostering acceptance and ensuring equitable access to these advanced technologies.

Future studies should explore patient attitudes toward AI in healthcare, particularly regarding their trust in AI-generated recommendations compared to traditional methods. Engaging patients in discussions about the benefits and limitations of AI in cancer detection will be essential for developing user-friendly and ethically sound systems.

The VIPER model demonstrates significant promise in the realm of cancer detection using genomic sequences. By integrating VIPER into clinical workflows, enhancing interactive decision support systems, incorporating multi-omics data, and improving model interpretability, the potential for innovative advancements in cancer diagnosis and treatment is substantial. However, as with any emerging technology, careful attention to ethical considerations and patient acceptance will be paramount to its successful implementation in healthcare.

VIII. CONCLUSION AND FUTURE WORK

In conclusion, the extended analysis of the results underscores the effectiveness of the VIPER model in the domain of cancer detection from genomic sequences. By providing a detailed explanation of each result metric and offering a thorough discussion of the implications, this section highlights the significant advancements achieved by VIPER. The model's ability to accurately detect cancer-causing mutations, evidenced by a training accuracy of 96.84%, a validation accuracy of 97.30%, and a training loss of 0.1548, sets a new benchmark in the field of genomic data analysis. The insights gained from this study pave the way for future research and development, with the potential to revolutionize mutation detection using artificial intelligence. Cancer detection is a significant challenge in bioinformatics. VIPER succeeds in overcoming obstacles such as identifying long-range dependencies in genomic data and optimizing computational efficiency through the use of Conv1D layers and Mamba blocks. These components contribute to its outstanding precision of 97.52%, recall of 96.51%, and F1 score of 97.13%.

While VIPER demonstrates strong performance, further hyperparameter tuning will be necessary when applied to more diverse and comprehensive datasets. Future research should consider optimizing the model architecture to accommodate variations in genomic sequences, ensuring VIPER remains computationally viable as the complexity of datasets increases. The success of the VIPER model in detecting cancer-causing mutations opens several avenues for future research. One potential direction is to explore its application to different types of mutations or cancer types, adapting the model to maintain the same level of accuracy and efficiency across varied genomic landscapes. Additionally, integrating other data modalities, such as clinical outcomes or multi-omics data, could further enhance the model's capacity to generate insights into cancer biology. Another area for future work involves investigating novel architectures or ensemble methods to further improve performance, particularly when faced with more extensive and heterogeneous datasets.

REFERENCES

[1] K. Bera, K. A. Schalper, D. L. Rimm, V. Velcheti, and A. Madabhushi, "Artificial intelligence in digital pathology—new tools for diagnosis and precision oncology," *Nature reviews Clinical oncology*, vol. 16, no. 11, pp. 703–715, 2019.

[2] T. Le Scao, A. Fan, C. Akiki, E. Pavlick, S. Ilić, D. Hesslow, R. Castagné, A. S. Luccioni, F. Yvon, M. Gallé *et al.*, "Bloom: A 176b-parameter open-access multilingual language model," 2023.

[3] Z.-Q. Zhao, P. Zheng, S.-t. Xu, and X. Wu, "Object detection with deep learning: A review," *IEEE transactions on neural networks and learning systems*, vol. 30, no. 11, pp. 3212–3232, 2019.

[4] K. Christofyllakis, J. T. Bittenbring, L. Thurner, M. Ahlgrimm, S. Stilgenbauer, M. Bewarder, and D. Kaddu-Mulindwa, "Cost-effectiveness of precision cancer medicine-current challenges in the use of next generation sequencing for comprehensive tumour genomic profiling and the role of clinical utility frameworks," *Molecular and Clinical Oncology*, vol. 16, no. 1, p. 21, 2022.

[5] A. Gu and T. Dao, "Mamba: Linear-time sequence modeling with selective state spaces," *arXiv preprint arXiv:2312.00752*, 2023.
[6] J. P. Amorim, P. H. Abreu, A. Fernández, M. Reyes, J. Santos, and M. H. Abreu, "Interpreting deep machine learning models: an easy guide for oncologists," *IEEE Reviews in Biomedical Engineering*, vol. 16, pp. 192–207, 2021.
[7] Y. Zhao, Z. Pan, S. Namburi, A. Pattison, A. Posner, S. Balachander, C. A. Paisie, H. V. Reddi, J. Rueter, A. J. Gill *et al.*, "Cup-ai-dx: A tool for inferring cancer tissue of origin and molecular subtype using rna gene-expression data and artificial intelligence," *EBioMedicine*, vol. 61, 2020.
[8] K. Sienkiewicz, J. Chen, A. Chatrath, J. T. Lawson, N. C. Sheffield, L. Zhang, and A. Ratan, "Detecting molecular subtypes from multi-omics datasets using sumo," *Cell reports methods*, vol. 2, no. 1, 2022.
[9] M. Mostavi, Y.-C. Chiu, Y. Huang, and Y. Chen, "Convolutional neural network models for cancer type prediction based on gene expression," *BMC medical genomics*, vol. 13, pp. 1–13, 2020.
[10] A. Cruz-Roa, H. Gilmore, A. Basavanhally, M. Feldman, S. Ganesan, N. N. Shih, J. Tomaszewski, F. A. González, and A. Madabhushi, "Accurate and reproducible invasive breast cancer detection in whole-slide images: A deep learning approach for quantifying tumor extent," *Scientific reports*, vol. 7, no. 1, p. 46450, 2017.
[11] P. Stanitsas, A. Cherian, A. Truskinovsky, V. Morellas, and N. Papanikolopoulos, "Active convolutional neural networks for cancerous tissue recognition," in *2017 IEEE International Conference on Image Processing (ICIP)*, 2017, pp. 1367–1371.
[12] A. Gupta and S. Sharma, "Genome sequence identification using neural network for breast cancer diagnosis," in *2023 3rd Asian Conference on Innovation in Technology (ASIANCON)*. IEEE, 2023, pp. 1–6.
[13] A. Suresh, R. R. Nair, E. Neeba, and S. A. Kumar, "Recurrent neural network for genome sequencing for personalized cancer treatment in precision healthcare," *Neural Process Lett*, vol. 55, pp. 1–10, 2021.
[14] F. Pouladi, H. Salehinejad, and A. M. Gilani, "Recurrent neural networks for sequential phenotype prediction in genomics," in *2015 International Conference on Developments of E-Systems Engineering (DeSE)*. IEEE, 2015, pp. 225–230.
[15] Y. Tao, C. Cai, W. W. Cohen, and X. Lu, "From genome to phenome: Predicting multiple cancer phenotypes based on somatic genomic alterations via the genomic impact transformer," in *PACIFIC SYMPOSIUM ON BIOCOMPUTING 2020*. World Scientific, 2019, pp. 79–90.
[16] A. Khan and B. Lee, "Gene transformer: Transformers for the gene expression-based classification of lung cancer subtypes," *arXiv preprint arXiv:2108.11833*, 2021.
[17] E. Nguyen, M. Poli, M. Faizi, A. Thomas, M. Wornow, C. Birch-Sykes, S. Massaroli, A. Patel, C. Rabideau, Y. Bengio *et al.*, "Hyenadna: Long-range genomic sequence modeling at single nucleotide resolution," *Advances in neural information processing systems*, vol. 36, 2024.
[18] C. Angermueller, T. Pärnamaa, L. Parts, and O. Stegle, "Deep learning for computational biology," *Molecular systems biology*, vol. 12, no. 7, p. 878, 2016.
[19] Y. Yuan, Y. Shi, C. Li, J. Kim, W. Cai, Z. Han, and D. D. Feng, "Deepgene: an advanced cancer type classifier based on deep learning and somatic point mutations," *BMC bioinformatics*, vol. 17, pp. 243–256, 2016.
[20] Y. Yuan, Y. Shi, X. Su, X. Zou, Q. Luo, D. D. Feng, W. Cai, and Z.-G. Han, "Cancer type prediction based on copy number aberration and chromatin 3d structure with convolutional neural networks," *BMC genomics*, vol. 19, pp. 1–8, 2018.
[21] J. Tian, M. Zhu, Z. Ren, Q. Zhao, P. Wang, C. K. He, M. Zhang, X. Peng, B. Wu, R. Feng *et al.*, "Deep learning algorithm reveals two prognostic subtypes in patients with gliomas," *BMC bioinformatics*, vol. 23, no. 1, p. 417, 2022.
[22] K. Krishnamachari, D. Lu, A. Swift-Scott, A. Yeraliyev, K. Lee, W. Huang, S. N. Leng, and A. J. Skanderup, "Accurate somatic variant detection using weakly supervised deep learning," *Nature Communications*, vol. 13, no. 1, p. 4248, 2022.
[23] R. J. Chen, C. Chen, Y. Li, T. Y. Chen, A. D. Trister, R. G. Krishnan, and F. Mahmood, "Scaling vision transformers to gigapixel images via hierarchical self-supervised learning," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2022, pp. 16 144–16 155.
[24] T.-H. Zhang, M. M. Hasib, Y.-C. Chiu, Z.-F. Han, Y.-F. Jin, M. Flores, Y. Chen, and Y. Huang, "Transformer for gene expression modeling (t-gem): an interpretable deep learning model for gene expression-based phenotype predictions," *Cancers*, vol. 14, no. 19, p. 4763, 2022.
[25] Y. Yuan, Y. Shi, X. Su, X. Zou, Q. Luo, D. D. Feng, W. Cai, and Z.-G. Han, "Cancer type prediction based on copy number aberration and chromatin 3d structure with convolutional neural networks," *BMC genomics*, vol. 19, pp. 1–8, 2018.
[26] S. Hafezqorani, "Machine learning in transcriptome analysis using long rna sequencing data," Ph.D. dissertation, University of British Columbia, 2024.
[27] S. M. E. Sahraeian, L. T. Fang, K. Karagiannis, M. Moos, S. Smith, L. Santana-Quintero, C. Xiao, M. Colgan, H. Hong, M. Mohiyuddin *et al.*, "Achieving robust somatic mutation detection with deep learning models derived from reference data sets of a cancer sample," *Genome Biology*, vol. 23, no. 1, p. 12, 2022.
[28] F. Martínez-Jiménez, F. Muiños, I. Sentís, J. Deu-Pons, I. Reyes-Salazar, C. Arnedo-Pac, L. Mularoni, O. Pich, J. Bonet, H. Kranas *et al.*, "A compendium of mutational cancer driver genes," *Nature Reviews Cancer*, vol. 20, no. 10, pp. 555–572, 2020.
[29] Y. Zhao, Z. Pan, S. Namburi, A. Pattison, A. Posner, S. Balachander, C. A. Paisie, H. V. Reddi, J. Rueter, A. J. Gill *et al.*, "Cup-ai-dx: A tool for inferring cancer tissue of origin and molecular subtype using rna gene-expression data and artificial intelligence," *EBioMedicine*, vol. 61, 2020.
[30] A. Gu, K. Goel, and C. Ré, "Efficiently modeling long sequences with structured state spaces," *arXiv preprint arXiv:2111.00396*, 2021.
[31] J. X. Chen, E. P. Wileyto, and M. C. Soulen, "Randomized embolization trial for neuroendocrine tumor metastases to the liver (retnet): study protocol for a randomized controlled trial," *Trials*, vol. 19, pp. 1–8, 2018.
[32] J. M. Johnson and T. M. Khoshgoftaar, "Survey on deep learning with class imbalance," *Journal of big data*, vol. 6, no. 1, pp. 1–54, 2019.
[33] S. Bai, J. Z. Kolter, and V. Koltun, "An empirical evaluation of generic convolutional and recurrent networks for sequence modeling. arxiv," *arXiv preprint arXiv:1803.01271*, vol. 10, 2018.

AMIT KUMAR BAIRWA serves as a distinguished academician in the School of Computer Science and Engineering at Manipal University Jaipur. He holds an impressive academic background, boasting qualifications such as B.E., M.Tech., and Ph.D. in Computer Science and Engineering. His expertise lies in the fields of Optimization, MANET, Network Security, Machine Learning. As an accomplished author, Dr. Bairwa has published 4 books and is actively involved in research across various areas of his interest. He is a valued member of prestigious organizations including ISTE, ACM, and IEEE, which reflects his commitment to professional growth and networking. In addition to his academic pursuits, Dr. Amit is highly proficient in handling real-time challenges related to network and Linux-based systems. He has earned several international certifications, such as CCNA, RHCE, and CEH, demonstrating his dedication to staying current with the latest industry practices. Dr. Amit also designed 4 grants and 8 copyrights. Furthermore, Dr. Bairwa has made significant contributions to the academic ecosystem during his tenure. He has taken initiatives to enhance the learning environment within the classroom, developed cutting-edge academic and professional programs, and adapted them to meet the evolving needs of the industry in today's dynamic business landscape. He has been awarded with so many titles in the field of teaching and research. His approach involves a well-balanced combination of theoretical concepts and industry-oriented training modules. Collaborating with fellow academicians, Dr. Amit Kumar Bairwa has played an integral role in setting and maintaining high curriculum standards. He actively contributes to the development of mission statements and establishes performance goals and objectives, ensuring that the educational institution consistently excels in its academic pursuits. He also organizes International Conference on Cyber Warfare, Security & Space Computing (SpacSec) and International Conference on Computation of Artificial Intelligence & Machine Learning (ICCAIML) on regular basis. Apart from this He is actively conducting workshops, expert talks and gave judgement to top events at university level.

SIDDHANTH BHAT is a dedicated scholar pursuing a Bachelor's in Computer Science with honours in Artificial Intelligence and Machine Learning at Manipal University Jaipur. His areas of expertise include Deep Learning, Natural Language Processing (NLP), and Data Science. Siddhanth has contributed to several research papers, with publications in IEEE journals and conferences, focusing on topics such as Image Recognition, Language Processing, Generative AI, and Data Analysis.

Professionally, Siddhanth serves as the Research Lead and Co-Founder at Xneuronz, and has interned at MuSigma, where he gained valuable experience in data analysis. He has also held various leadership roles, including President of Panacea, the Computer Communications Society, and General Secretary of the Artificial Intelligence Department Community. His active participation in academic events includes contributions to the organizing committees for the International Conference on Computation of Artificial Intelligence Machine Learning (ICCAIML) and the International Conference on Innovations in Computational Intelligence and Computer Vision (ICICV).

In addition to his academic and professional pursuits, Siddhanth holds certifications in TensorFlow, AWS, and Oracle Database Foundations, reflecting his commitment to continuous learning. Focused on applying AI to solve real-world challenges, he is passionate about advancing the field of computer science through both innovative research and practical applications.

TANISHK SAWANT is an aspiring IT professional currently pursuing a B.Tech. in Information Technology at Manipal University Jaipur. With a strong interest in Natural Language Processing (NLP) and Neural Networks, Tanishk has developed expertise in these domains, focusing on their real-world applications. He is currently interning at Fugumobile in Mumbai, where he is working on cutting-edge projects in NLP, Speech Synthesis, Voice Cloning, and Computer Vision. This role allows him to apply his knowledge in creating innovative solutions, further enhancing his technical skills and practical experience. Previously, Tanishk interned at MuSigma, gaining valuable experience in data analysis. His academic and professional experiences have equipped him with a solid foundation in data processing and machine learning techniques. Driven by a passion for exploring the potential of AI and machine learning, Tanishk is dedicated to advancing his knowledge and contributing to the development of technology that addresses real-world challenges.

PRABHATH VARMA is an aspiring Computer Science engineer currently pursuing his Bachelor of Technology at Manipal University Jaipur. His academic interests span AI, Data Science, and Network Systems, with a growing expertise in areas like Microsoft AI Fundamentals and Azure DevOps. Prabhath has gained hands-on experience during his internship at Trianz, where he worked on data cleaning, client presentations, and mock pitch sessions. Prabhath has contributed to multiple innovative projects, including research on human gait analysis, focusing on its biomechanical impacts and safety applications. He also explored Disaster response coordination platforms connecting social media with the nearest Disaster response team. Additionally, he orchestrated a Premier League prediction tournament, using advanced Excel techniques to automate scoring and provide real-time updates. Certified in Microsoft AI Fundamentals, IBM Data Science, and Oracle Database Foundations, Prabhath is dedicated to expanding his knowledge and skills in the field. Outside of academics, he is an accomplished speed skater, having represented India internationally and earned multiple national medals. His passion for sports extends to athletics, where he serves as the co-captain of his university's team.

MANOJ R is a distinguished academician with a solid educational foundation, holding a B.E. and an MTech. He is currently pursuing a PhD in Computer Science and Engineering with a specialization in blockchain technology. He serves as an assistant professor in the Department of Computer Science & Engineering at Manipal Institute of Technology, part of the Manipal Academy of Higher Education (MAHE), located in Manipal, Karnataka, India. With over 15 years of academic experience, Manoj has established himself as a prolific contributor to his field. He has presented and published more than 20 papers in prestigious international journals and conferences, demonstrating his commitment to advancing knowledge and research. Additionally, he has provided his expertise as a reviewer for numerous research papers submitted to reputable journals and conferences. Manoj's research interests are broad and impactful, encompassing cryptography and network security, as well as blockchain technology and its various applications. His work in these areas underscores his dedication to exploring and addressing critical issues in computer science and engineering.